

Architecture: Microservices

17-313 Fall 2025

Foundations of Software Engineering

<https://cmu-17313q.github.io>

Eduardo Feo Flushing

Learning Goals

- Contrast the monolithic application design with a modular design based on microservices.
- Reason about tradeoffs of microservices architectures.
- Principles of microservices: how to benefit and avoid their pitfalls

Outline


- From Monoliths to Service Oriented Architecture
 - Case Study: Chrome Web Browser
- Microservices
 - Monolith vs Microservices
 - Advantages
 - Challenges
- Microservices: Principles
- Serverless

Before we get to
microservices...



MONOLITHS



Monolithic



Monolithic styles




Source: <https://www.seobility.net> (CC BY-SA 4.0)

Modularity comes in many ways

- **Plug-in architectures**

- Distinct code repositories, joined to a monolithic run-time
- Examples: Linux kernel modules, NodeBB themes, VS Code extensions
- Separates development, but runs as “one”




Modularity comes in many ways

- Plug-in architectures
- Service-oriented architectures
 - Distinct processes communicating via messages (e.g., Web browsers)
 - Separates run-time resource management and failure / security issues.
- Distributed microservices
 - Independent, autonomous services communicating via web APIs
 - Separates almost all concerns


Separation of concerns

SERVICE-ORIENTED ARCHITECTURE

Monolithic




Service-oriented





Example: Chrome

Web Browsers




Source: <https://developers.google.com/web/updates/2018/09/inside-browser-part1> (CC BY 4.0)

The evolution of browser architectures




Source: <https://developers.google.com/web/updates/2018/09/inside-browser-part1> (CC BY 4.0)

Service-oriented browser architecture





Source: <https://developers.google.com/web/updates/2018/09/inside-browser-part1> (CC BY 4.0)

Implementation: Multi-process browser with IPC




Source: <https://developers.google.com/web/updates/2018/09/inside-browser-part1> (CC BY 4.0)

Benefits of a service-oriented browser architecture




Source: <https://developers.google.com/web/updates/2018/09/inside-browser-part1> (CC BY 4.0)

Navigating to a web site uses service requests




Source: https://developers.google.com/web/updates/2018/09/inside_browser_part1 (CC BY 4.0)

Navigating to a web site uses service requests




Source: https://developers.google.com/web/updates/2018/09/inside_browser_part1 (CC BY 4.0)

Navigating to a web site uses service requests




Source: <https://developers.google.com/web/updates/2018/09/inside-browser-part1> (CC BY 4.0)

Navigating to a web site uses service requests




Source: <https://developers.google.com/web/updates/2018/09/inside-browser-part1> (CC BY 4.0)




Navigating to a web site uses service requests



Navigating to a web site uses service requests



Source: https://developers.google.com/web/updates/2018/09/inside_browser_part1 (CC BY 4.0)



<https://webperf.tips/tip/browser-process-model/>

One browser, many processes

```
547187 ? SNI 0:00 /opt/google/chrome/chrome_crashpad_handler --no-periodic-tasks --m  
--annotation=ver=128.0.6613.137 --initial-client-fd=4 --shared-client-connection  
547193 ? SN 0:00 /opt/google/chrome/chrome --type=zygote --no-zygote-sandbox --crash  
547194 ? SN 0:00 /opt/google/chrome/chrome --type=zygote --crashpad-handler-pid=547185  
547196 ? SN 0:04 /opt/google/chrome/chrome --type=zygote --crashpad-handler-pid=547185 --enab  
--enable-crash-reporter=, --crashpad-handler-pid=547185 --enable-crash-reporter=, --ch  
547220 ? SNI 48:08 /opt/google/chrome/chrome --type=gpu-process --crashpad-handler-pid=547185 --enable-crash-reporter:  
--shared-files --field-trial-handle=3,i,16924338087686289425,14108195033817998739,262144 --disable-features=EyeDropper --vai  
547222 ? SNI 21:09 /opt/google/chrome/chrome --type=utility --utility-sub-type=network.mojom.NetworkService --lang=en-US  
--disable-features=EyeDropper --variations-seed-version=20240923-050122.947000  
547224 ? SNI 0:14 /opt/google/chrome/chrome --type=utility --utility-sub-type=storage.mojom.StorageService --lang=en-US -  
--disable-features=EyeDropper --variations-seed-version=20240923-050122.947000  
547410 ? SNI 2:30 /opt/google/chrome/chrome --type=renderer --crashpad-handler-pid=547185 --enable-crash-reporter=, --r  
--shared-files=v8_context_snapshot_data:100 --field-trial-handle=3,i,16924338087686289425,14108195033817998739,262144 --di  
547533 ? SNI 1:23 /opt/google/chrome/chrome --type=renderer --crashpad-handler-pid=547185 --enable-crash-reporter=, --r  
--shared-files=v8_context_snapshot_data:100 --field-trial-handle=3,i,16924338087686289425,14108195033817998739,262144 --di  
547594 ? SNI 10:27 /opt/google/chrome/chrome --type=utility --utility-sub-type=audio.mojom.AudioService --lang=en-US --se  
--disable-features=EyeDropper --variations-seed-version=20240923-050122.947000  
547755 ? SNI 0:08 /opt/google/chrome/chrome --type=renderer --crashpad-handler-pid=547185 --enable-crash-reporter=, --r  
--shared-files=v8_context_snapshot_data:100 --field-trial-handle=3,i,16924338087686289425,14108195033817998739,262144 --di  
554653 ? SNI 2:43 /opt/google/chrome/chrome --type=renderer --crashpad-handler-pid=547185 --enable-crash-reporter=, --e  
--shared-files=v8_context_snapshot_data:100 --field-trial-handle=3,i,16924338087686289425,14108195033817998739,262144 --di  
554722 ? SNI 0:11 /opt/google/chrome/chrome --type=renderer --crashpad-handler-pid=547185 --enable-crash-reporter=, --n  
--shared-files=v8_context_snapshot_data:100 --field-trial-handle=3,i,16924338087686289425,14108195033817998739,262144 --di  
.....  
1500520 ? SNI 0:00 /opt/google/chrome/chrome --type=utility --utility-sub-type=data_decoder.mojom.DataDecoderService --la  
--field-trial-handle=3,i,16924338087686289425,14108195033817998739,262144 --disable-features=EyeDropper --variations-seed-  
1500532 ? SNI 0:00 /opt/google/chrome/chrome --type=renderer --crashpad-handler-pid=547185 --enable-crash-reporter=, --
```

Service-oriented architecture: More benefits


- Ability to change components independently
- Independent processes (Isolation, Security)
- Focusing on doing one thing well

MICROSERVICES


Why do we need microservices if we already have modular architectures and SOA?

- **Monoliths:**
 - Often lack modularity, and even “modular” monoliths are tightly coupled in runtime.
- **Service oriented architectures:**
 - More flexible, but often heavyweight and centralized.


Monolithic



Service oriented



Microservices





*“Small autonomous services that work well
together”*

Sam Newman

Microservices



COMCAST




UBER

GROUPON®

Netflix Microservices – App Boot



- Recommendations
- Trending Now
- Continue Watching
- My List
- Metrics



Microservices: Demo

ONLINEBOUTIQUE

\$ USD 

local




Hot Products



 Sunglasses \$19.99	 Tank Top \$18.99	 Watch \$109.99
 Loafers	 Hair Dryer	 Candle Holder

Service	Language	Description
frontend	Go	Exposes an HTTP server to serve the website. Does not require signup/login and generates session IDs for all users automatically.
cartservice	C#	Stores the items in the user's shopping cart in Redis and retrieves it.
productcatalogservice	Go	Provides the list of products from a JSON file and ability to search products and get individual products.
currencyservice	Node.js	Converts one money amount to another currency. Uses real values fetched from European Central Bank. It's the highest QPS service.
paymentservice	Node.js	Charges the given credit card info (mock) with the given amount and returns a transaction ID.
shippingservice	Go	Gives shipping cost estimates based on the shopping cart. Ships items to the given address (mock)
emailservice	Python	Sends users an order confirmation email (mock).
checkoutservice	Go	Retrieves user cart, prepares order and orchestrates the payment, shipping and the email notification.
recommendationservice	Python	Recommends other products based on what's given in the cart.
adservice	Java	Provides text ads based on given context words.
loadgenerator	Python/Locust	Continuously sends requests imitating realistic user shopping flows to the frontend.

Home Page




Checkout Screen



Services

 Reset  Refresh



Monoliths vs Microservices

Activity: In teams of 3-4

What are the consequences of this architecture? On:

- Scalability
- Reliability
- Performance
- Development
- Maintainability
- Testability
- Ownership




Scalability


A monolithic application puts all its functionality into a single process...




... and scales by replicating the monolith on multiple servers



A microservices architecture puts each element of functionality into a separate service...



... and scales by distributing these services across servers, replicating as needed.




Source: <http://martinfowler.com/articles/microservices.html>

Types of scaling: vertical vs. horizontal


Vertical Scaling

Increase or decrease the capacity of existing services/instances.





Horizontal Scaling

Add more resources like virtual machines to your system to spread out the workload across them.




Data management and consistency




Source: <http://martinfowler.com/articles/microservices.html>

Deployment and Evolution



monolith - multiple modules in the same process




microservices - modules running in different processes


Source: <http://martinfowler.com/articles/microservices.html>

MICROSERVICES: PRINCIPLES


Principles




Sam Newman's Principles of Microservices





Sam Newman's Principles of Microservices







Sam Newman's Principles of Microservices








Sam Newman's Principles of Microservices











Sam Newman's Principles of Microservices















Sam Newman

Principle 1: Domain-driven modeling


- Model services around business capabilities



Principle 1: Domain-driven modeling




Principle 1: Domain-driven modeling



Principle 2: Culture of Automation

- API-Driven Machine Provisioning
- Continuous Delivery
- Automated Testing



API-Driven Machine Provisioning

Example: Infrastructure as code (IaC)




Image source: <https://learn.microsoft.com/en-us/devops/deliver/what-is-infrastructure-as-code>

Continuous Delivery

More on this topic later






Image Source: <https://learn.microsoft.com/en-us/azure/architecture/microservices/ci-cd>


Principle 3: Hide implementation details

- Design carefully your APIs
- It's easier to expose some details later than hide them
- Do not share your database!




Principle 3: Hide implementation details

Recall: Encapsulation in OOP




Sharing database: Anti-pattern



Principle 4: Decentralized Governance

- Mind Conway's Law
- You Build It, You Run It
- Embrace team autonomy
- Internal Open Source Model







Mind Conway's Law



“Organizations which design systems are constrained to produce designs which are copies of the communication structures of these organizations”

- Melvin Conway (1967).

Mind Conway's Law



“Products” not “Projects”


YOU BUILD IT YOU RUN ~~AWAY~~ IT

"The traditional model is that you take your software to the wall that separates development and operations, and throw it over and then forget about it. Not at Amazon. You build it, you run it. This brings developers into contact with the day-to-day operation of their software. It also brings them into day-to-day contact with the customer. This customer feedback loop is essential for improving the quality of the service."


-- Werner Vogels in "A conversation with Werner Vogels" in ACM Queue, May 2006

Principle 5: Deploy Independently


- One Service Per OS
- Consumer-Driven Contracts
- Multiple co-existing versions




One Service Per OS




Q.
What is the problem
with this deployment?



Consumer-Driven Contracts




Multiple coexisting versions



Principle 6: Consumer First

- Encourage conversations
- API Documentation
- Service Discovery




Encourage conversations



VS





API Documentation




Swagger™


Principle 7: Isolate Failure

- Avoid cascading failures
- Timeouts between components
- **Fail fast** aka *Design for Failure*
 - Bulkheading / Circuit breakers





Closed circuit




Open circuit

Image source: blogs.halodoc.io


Principle 8: Highly Observable



- Standard Monitoring
- Health-Check Pages
- Log and Stats aggregation
- Downstream monitoring



Principle 8: Highly Observable

- Standard Monitoring
- Health-Check Pages
- Log and Stats aggregation
- Downstream monitoring





Sam Newman

Are microservices always the right choice?

Summary: Microservice challenges

- Too many choices
- Delay between investment and payback
- Complexities of distributed systems
 - network latency, faults, inconsistencies
 - testing challenges
- Monitoring is more complex
- More system states
- More points of failure
- Operational complexity
- Frequently adopted by breaking down a monolithic application



Summary: Advantages of Microservices

- Ship features faster and safer
- Scalability
- Target security concerns
- Allow the interplay of different systems and languages, no commitment to a single technology stack
- Easily deployable and replicable
- Embrace uncertainty, automation, and faults

Microservices overhead

for less-complex systems, the extra baggage required to manage microservices reduces productivity

